# SciComp Project 3
# Week 4+5 (10 points)

September 27, 2021
To be handed in October 11 before 12:00 noon

# 1 Background

## 1.1 Importance of Rapid Convergence

Finding a minimum energy configuration for a system of mutiple atoms is a widely applicaple for designing molecular systems in solid state physics, chemistry, medicine and other fields. Many applications will require tens to hundreds of atoms in the model. Calculating the energy involves solving the many-body Schrødinger equation, in which all the electrons interact with each other.

With hundreds of atoms and thousands of electrons, solving the many-body Schrödinger equation (even approximately) can take weeks of CPU-time to calculate a single energy point, whereby a full optimization easily runs over many years of CPU-time. Even on a parallel computer with e.g. 1000 CPU cores, optimizing the geometry of a single molecule can easily take months in "human time".

When every evaluation of the energy function takes, for example, a CPU-week to compute (burning large amounts of $CO_2$), the number of function evaluations matter: a problem common to many fields, not just quantum chemistry.

In this assignment, you will try your hand at optimizing a simple system of interacting atoms. However, we do not have months to wait for computations to complete, so we will use an extremely simple approximation to the potential energy, the *Lennard-Jones*-potential. This approximation is appropriate for noble gas atoms, yielding a crude picture of the formation; yet we will pretend that it is a full *ab initio* quantum chemical energy calculation, and assume that it takes a week to complete. Thus, you will be asked to *count energy function evaluations* and report the time spent.

## 1.2 Minimizing the potential in a cloud of Argon atoms using the Lennard-Jones potential

The Lennard-Jones potential works for systems of neutral atoms or molecules. It assumes that no new electronic bonds nor new molecules are formed, and that the potential can be modeled only with a short-range repulsion force (arising from the Pauli exclusion principle) and a long-range van der Waals attraction. This yields a classical description of the system, in which the neutral atoms or molecules move as classical particles, interacting only through this simple potential.

The Lennard-Jones potential can be written in several ways, but the most common is:

$$v_{LJ}(r_{ij}) = 4\epsilon \left( \left( \frac{\sigma}{r_{ij}} \right)^{12} - \left( \frac{\sigma}{r_{ij}} \right)^{6} \right) \tag{1}$$

$v_{LJ}(r_{ij})$ is the potential at distance $r_{ij} = \|\mathbf{x}_i - \mathbf{x}_j\|_2$ between two atoms $i$ and $j$. The repulsive force is stronger than the van der Waals-force, but decreases more rapidly with distance. Thus, for a pair of atoms, there is a distance where the total potential is minimal, the potential well. $\epsilon$ is this minimal potential between two atoms, and $\sigma$ is the inter-particle distance where the potential is zero: $v_{LJ}(\sigma) = 0$. The constants $\epsilon$ and $\sigma$ are generally found by fitting to proper ab initio calculations, or using experimental data.

The full potential LJ-energy of $N$ neutral particles is the sum of pair-potentials:

$$V_{LJ}(\mathbf{X}) = \sum_{i=1}^{N} \sum_{j=i+1}^{N} v_{LJ}(r_{ij}) \tag{2}$$

where $\mathbf{X} \in \mathbb{R}^{N \times 3}$ is the matrix of coordinates for the $N$ particles.

In this assignment, we will be modeling argon, for which

$$\sigma = 3.401\text{Å and } \epsilon = 0.997\text{kJ/mol} \tag{3}$$

## 1.3 Array Programming and Provided Functions

Try to use array programming when programming with Python/Numpy or Matlab, as the speed difference can be up to a factor of hundreds. As an example of how to use array programming, here is how a distance matrix (to be used in the LJ-potential) can be calculated with array operations using NA = newaxis (as I explain further here):

$$D_{ij} = \|\mathbf{x}_i - \mathbf{x}_j\|_2 = \|\mathbf{x}_{i\mathbf{j}} - \mathbf{x}_{ij}\|_2 \tag{4}$$

In Eq. (4), $\mathbf{x}_i$ is extended with a newaxis $\mathbf{j}$ and $\mathbf{x}_j$ with a newaxis $\mathbf{i}$ to put them on the same footing, so that the pair differences can be calculated as elementwise subtraction $\mathbf{x}_{i\mathbf{j}} - \mathbf{x}_{ij}$. In Numpy code, it looks as follows:

```python
# points:          (N,3)-array of (x,y,z) coordinates for N points
# distance(points): returns (N,N)-array of inter-point distances:
def distance( points ):
        displacement = points[:,NA] - points[NA,:]
        return sqrt( sum(displacement*displacement, axis=-1) )
```

In the file LJhelperfunctions.py we have collected the functions you will need to calculate the Lennard-Jones potential and its gradient. The function V=LJ(sigma, epsilon) takes two scalar arguments: sigma and epsilon, and returns a function V, which calculates the Lennard-Jones potential for the supplied values of $\sigma$ and $\epsilon$. This type of programming construct is called *closure*, and is also useful in restricting multidimensional functions along a line through space (as you will see when implementing line searches). We also provide a function V(points), generated from LJ with the experimental values for Argon particles, specified in equation 3. The function takes a $N \times 3$ array of atomic positions and returns the strength of the Lennard-Jones potential between particles at these positions. The implementation is as follows

```python
def LJ(sigma, epsilon):
    def V(points):                 # points: (N,3)-array of (x,y,z) coordinates for N points
        dist = distance(points)
        fill_diagonal(dist, 1)             # Fill diagonal with 1 to avoid division by zero
        f = sigma/dist                     # Dimensionless reciprocal distance
        pot = 4*epsilon*(f**12 - f**6)   # Calculate the interatomic potentials
        fill_diagonal(pot, 0)              # Remove diagonal (no self-interaction)
        return sum(pot)/2
    return V
```

# 2 Questions for Week 4: Solving Nonlinear Equations

## 2.1 Solving Nonlinear Equations in 1D

a. Using the provided function `V`, 1) plot the strength of the potential between two particles $\mathbf{x}_0 = (x, 0, 0)$ and $\mathbf{x}_1 = (0, 0, 0)$. Do this by writing a function that takes in a single scalar argument $x$, and calculates the potential using `V`, with $x$ ranging between 3 and 11. 2) Plot the strength of the potential between four particles: $\mathbf{x}_0$ and $\mathbf{x}_1$ as before, $\mathbf{x}_2 = (14, 0, 0)$ and $\mathbf{x}_3 = (7, 3.2, 0)$, again with $x$ ranging from 3 to 11. Do this by again writing a function that takes a single scalar argument $x$ and returns the potential strength.

b. Write a bisection root finding function `x, n_calls = bisection_root(f,a,b,tolerance=1e−13)` that finds $\mathbf{x}$ such that $f(\mathbf{x}) = 0$ given a bracket $x \in [a; b]$, and counts the number of calls to the function $f$. (A reasonable length is 8-12 lines of code).

In this assignment, let the convergence test be on how close we get $f(x)$ to zero. Test it to find the zero of the the LJ-potential between two argon atoms as a function of interatomic distance, and verify that you get $x = \sigma$. How many calls to the energy function were needed to get from the start bracket $[a, b] = [2, 6]$ to $|f(x)| < 10^{-13}$?

c. The derivative of the pair-potential $v_{LJ}(r) = 4\epsilon((\sigma/r)^{12} - (\sigma/r)^6)$ is

$$\frac{d}{dr} v_{LJ}(r) = 4\epsilon \left( \frac{6\sigma^6}{r^7} - \frac{12\sigma^{12}}{r^{13}} \right)$$

Write a Newton-Rhapson solver `x, n_calls = newton_root(f,df,x0,tolerance,max_iterations)`, (a reasonable length is 4-8 lines of code), and test it in the same way as above. For simplicity, assume that a call to the derivative `df` has the same cost as a call to `f`. How many calls were needed to get from $x_0 = 2$ to $|f(x^*)| < 10^{-12}$, i.e., 12 decimal digits for $x^*$ after the comma?

d. Make a combination of Newton-Rhapson and bisection that is *guaranteed to converge* (when run on an interval bracketing a zero), but takes advantage of the quadratic convergence of Newton-Rhapson iteration, and test it on the same example. How many calls to the LJ-energy function was needed to get from $x_0 = 2$, $[a, b] = [2, 6]$ to obtain $|f(x^*)| < 10^{-13}$?

**Note:** *If you have trouble completing this step, simply skip it and move on to the remaining questions, which only requires your bisection root solver to work. You can always return to solve it once you have completed tasks (e) and (f).*

## 2.2 Solving $N$-dimensional Nonlinear Equations

Using the chain rule for derivatives $f(g(x))' = f'(g(x))g'(x)$, we can write down the derivative of a pair-potential with respect to the position of one of the particles:[1]

$$\frac{\partial}{\partial \mathbf{x}_i} v_{LJ}(r_{ij}) = 4\epsilon \left( \frac{6\sigma^6}{r_{ij}^7} - \frac{12\sigma^{12}}{r_{ij}^{13}} \right) \frac{\partial r_{ij}}{\partial \mathbf{x}_i}$$

$$= 4\epsilon \left( \frac{6\sigma^6}{r_{ij}^7} - \frac{12\sigma^{12}}{r_{ij}^{13}} \right) \frac{\mathbf{x}_i - \mathbf{x}_j}{r_{ij}}$$

(5)

---

[1]NB: You don't need to understand these derivations to solve the problem, as the actual code for the gradient is provided to you. The derivations are here to show you how to do it yourself in the future.

This we can then use to find an expression for the gradient of the total LJ-energy:

$$\frac{\partial}{\partial \mathbf{x}_k} V_{LJ}(\mathbf{X}) = \frac{\partial}{\partial \mathbf{x}_k} \frac{1}{2} \sum_{i=1}^{N} \sum_{j \neq i} v_{LJ}(r_{ij})$$

$$= \frac{\partial}{\partial \mathbf{x}_k} \sum_{i=1}^{N} \frac{1}{2} \sum_{j \neq i}^{N} (\delta_{i,k} + \delta_{j,k}) v_{LJ}(r_{ij})$$

$$= \frac{\partial}{\partial \mathbf{x}_k} \sum_{j \neq k}^{N} v_{LJ}(r_{kj}) \tag{6}$$

$$= 4\epsilon \sum_{j \neq k}^{N} \left( \frac{6\sigma^6}{r_{kj}^7} - \frac{12\sigma^{12}}{r_{kj}^{13}} \right) \frac{\mathbf{x}_k - \mathbf{x}_j}{r_{kj}}$$

The position of all the particles is a point $\mathbf{X} \in \mathbb{R}^{3N}$, which we organize as a $(N,3)$-array, so that $\mathbf{x}_i = (x_i, y_i, z_i)$. The total potential energy is a function $V_{LJ} : \mathbb{R}^{3N} \to \mathbb{R}$, and the gradient taken at any particular configuration of particle positions is hence a $3N$-dimensional vector: $V_{LJ}(\mathbf{X}) \in \mathbb{R}^{3N}$. Thus, the gradient to $V_{LJ}$ is a function $\nabla V_{LJ} : \mathbb{R}^{3N} \to \mathbb{R}^{3N}$. The negative of the gradient is the *force* acting on the system, and its direction is that in which the potential decreases most rapidly. But notice that it is a $3N$-dimensional direction: it acts on *all* the particles at once.

An implementation of the gradient is given in the LJhelperfunctions.py file as well, again using closure. The function LJgradient generates a function gradV, which in turn allows you to calculate the gradient of the system efficiently. The implementation is as follows:

```
def LJgradient(sigma, epsilon):
    def gradV(X):
        d = X[:, NA] - X[NA, :]
        r = sqrt(sum(d*d, axis=-1))

        fill_diagonal(r, 1)

        T = 6*(sigma**6)*(r**-7)-12*(sigma**12)*(r**-13)
        # (N,N)-matrix of r-derivatives
        # Using the chain rule , we turn the (N,N)-matrix of r-derivatives into
        # the (N,3)-array of derivatives to Cartesian coordinate: the gradient.
        # (Automatically sets diagonal to (0,0,0) = X[ i]-X[ i ])
        u = d/r[:, :, NA]
        # u is (N,N,3)-array of unit vectors in direction of X[ i ]-X[ j ]
        return 4*epsilon*sum(T[:, :, NA]*u, axis=1)
    return gradV
```

Again, a function gradV is provided, which uses the experimental values for Argon.

Finding the minima in $\mathbb{R}^{3N}$ of the potential involves finding points where its $3N$-dimensional gradient is $\mathbf{0}$. An important component needed in next week's work is called *line search*, where we search for a point along a single direction $\mathbf{d} \in \mathbb{R}^{3N}$ at which the gradient *along this line* is zero. That is, we want to start in a point $\mathbf{X}_0 \in \mathbb{R}^{3N}$, and then find $\alpha \in [0; b]$ such that $0 = \mathbf{d} \cdot \nabla V_{LJ}(\mathbf{X}_0 + \alpha \mathbf{d})$, i.e., the gradient has no component in the direction of $\mathbf{d}$. This finds an optimum for the one-dimensional function $V_{LJ}(\mathbf{X}_0 + \alpha \mathbf{d})$.

For a visual guide to how line searches work, look at the Jupyter notebook on the subject, available here, and in PDF form here. It goes over how to restrict a multidimensional function onto a line through space using closure, such that it on a scalar input, and how this applies to line searches.

e. Look at the gradient of the 2-particle system in question (a) with $\mathbf{x}_1 = (0, 0, 0)$ and $\mathbf{x}_0 = (x, 0, 0)$, $x \in [3; 10]$. Why are exactly two components nonzero? Why are they equal and opposite? Plot the nonzero component for the derivative of the $x$-coordinate of $\mathbf{x}_0$ $(0, 0$-coordinate of gradient) together with the potential, and notice the relationship between the zero of the derivative and the minimum of the potential.

Next look at the gradient for the 4-particle system from (a) at one of the minima of your plot. Why is the gradient not zero?

f. Write a function alpha, ncalls = **linesearch**(F,X0, d, alpha_max, tolerance, max_iterations) that takes a function $\mathbf{F} : \mathbb{R}^{N \times 3} \to \mathbb{R}^{N \times 3}$, a start position $\mathbf{X}_0 \in \mathbb{R}^{N \times 3}$ and finds the zero along the line-segment $\mathbf{X}_0 + \alpha \mathbf{d}$ of $\mathbf{d} \cdot \mathbf{F}(\mathbf{X}_0 + \alpha \mathbf{d})$.[2] Use your bisection solver at this point, as using Newton-Rhapson requires second derivatives when $\mathbf{F}$ is the gradient.

Test your function by finding the minimum along $\mathbf{X}_0 + \alpha \mathbf{d}$ with

$$\mathbf{X}_0 = \begin{bmatrix} 4 & 0 & 0 \\ 0 & 0 & 0 \\ 14 & 0 & 0 \\ 7 & 3.2 & 0 \end{bmatrix}$$

and $\mathbf{d} = -\nabla V_{LJ}(\mathbf{X}_0)$, and with $\alpha \in [0; 1]$.

# 3 Questions for Week 5: Nonlinear Optimization

While it is possible to find a minimum by solving the equation $\nabla V_{LJ}(\mathbf{X}) = \mathbf{0}$, simply using the root solver has two disadvantages: First, for 1D line-searches, we can evaluate $f(\mathbf{X} + \alpha \mathbf{d})$ instead $\nabla f(\mathbf{X} + \alpha \mathbf{d}) \cdot \mathbf{d}$, usually costing significantly less per call; Second, in order to make sure that we find a minimum and not a maximum or saddle point, we need to ensure that the Hessian is positive definite. This can be built into the computation. Hence, it is better to use specialized optimization functions.

In this week's questions, you'll program optimization functions and apply it to the $N$-body potential $V_{LR}(\mathbf{X})$ to find minimal-energy configurations. I suggest that you program your optimizers in a general way, such that they take as input a flat vector (in the case of $N$ particles, a flat $3N$-dimensional vector X0).

In python, you can derive a function $\tilde{f} : \mathbb{R}^{3N} \to \mathbb{R}$ from $f : \mathbb{R}^{N \times 3} \to \mathbb{R}$ and $\nabla f : \mathbb{R}^{N \times 3} \to \mathbb{R}^{N \times 3}$ to $\nabla \tilde{f} : \mathbb{R}^{3N} \to \mathbb{R}^{3N}$ with the wrappers:

```
def flatten_function(f):
  return lambda X: f(X.reshape(-1,3)) # Reshape with '-1' means "whatever is left".


def flatten_gradient(f):
  return lambda X: f(X.reshape(-1,3)).reshape(-1)
```

In the provided functions we include both this wrapper and a flattened version of the potential (flat_V) and gradient (flat_gradV). With Matlab, you can use '@(X)'.

## Geometric Intuition for the Lennard-Jones Potential

It can be hard to figure out what to expect from the optimal particle configurations. Fortunately we can use some physics and a bit of geometric intuition to arrive at a good guess for the optimal configurations (and how to spot these).

---

[2]You can use either a closure or a lambda expression to define the one-dimensional restriction of $\mathbf{F}$ to the line.

To help with this, we have put together the following Jupyter notebook on the subject, available here, and in PDF form here. Do not worry if you do not quite grasp the code, the actual geometric thinking is more important.

## Questions Building on Monday's Lecture

Download the data file ArStart.npz (for Python) or ArStart.mat (for Matlab), which defines randomly generated starting points for $N = 2, 3, \ldots, 9$.

g. Write a function x_opt, n_calls = **golden_section_min**(f,a,b,tolerance=1e−3) that finds the minimum of a 1D-function on a unimodal interval $x \in [a; b]$, and use it to obtain the same $\alpha$ as you did in (f), but without using the gradient. Next, use it to obtain the optimal (minimal-energy) distance $r_0$ between two Ar atoms.

h. Write a function X_opt, N_calls, converged = **BFGS**(f,gradf, X,tolerance=1e−6, max_iterations=10000) which implements BFGS either with the direct step (approximating the Hessian) or inverse steps (approximating the Hessian inverse), as covered in the lectures. converged is a boolean which is True iff you converged to the requested tolerance within max_iterations steps.

Test it to find the minimum of the two-particle system, but using the N-particle potential for the two points defined by the array Xstart2 in ArStarts. Confirm that this yields the same $r_0$ as above.

i. Apply your BFGS-minimizer to the starting geometries from ArStarts, starting with $N = 2$ and stopping when you reach an $N$ you can't get to converge. Inspect the distance matrix. How many distances are within 1% of the two-particle optimum $r_0$? You can count them automatically by writing **sum**(**abs**(D−r0)/r0 <= 0.01) if D is the distance matrix.

Inspect the results in 3D[3]. The true minimum corresponds to the zero-temperature configuration of the system if left to slowly cool. A good optimum should look something like a single lattice, where as many atoms as possible are trapped in the potential well of one or more neighbours. Show the 3D picture for $N = 3$ and your highest converged $N$.

j. Add a line-search step to your BFGS-function, such that your change in $\mathbf{x}$ becomes $\alpha \Delta \mathbf{x}$ instead of directly using the $\Delta \mathbf{x}$ obtained from $\mathbf{B}_{k+1} \Delta \mathbf{x}_k = -\nabla f(\mathbf{X}_k)$ (or $\Delta \mathbf{x}_k = -\tilde{\mathbf{B}}_{k+1} \nabla f(\mathbf{X}_k)$ for inverse BFGS iteration). Use both to find the minima for the N-particle starting positions from ArStarts for increasing $N$. Does it affect the number of steps needed to converge? Can you get higher $N$ to converge than before? Does it change the quality of the solution (the number of van der Waals "bonds" between atoms). What about the total number of function calls needed to converge?

Finally, a bonus task. You can leave this out and still obtain a full score, or you can solve it and get extra points towards a full score: to make up for not completing previous tasks, or for your own satisfaction.

k. Implement one of the meta-heuristics of your own choice, and use it to minimize the energy for $N = 5$, $N = 9$ and $N = 20$. Are you now able to obtain a good solution for higher $N$? Show the 3D picture and the number of van der Waals "bonds".

Compare the convergence to that of the dampened BFGS from (j): Trace the energies and accumulated number of function calls at each iteration step, and plot the energy 1) per iteration step, and 2) as a function of accumulated number of function calls.

---

[3]In Matlab you can use scatter3, in Python you can use e.g. matplotlib's Axes3D and ax.scatter